

# Worst-case Optimal Join in Clojure

Finn Völkel

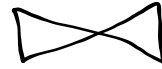
26.10.2022

# Joins

R(a,c)



T(a,b)

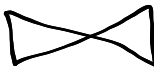
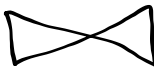


S(b,c)

```
SELECT
  *
FROM
  R, T, S
WHERE
  R.a = T.a AND T.b = S.b AND R.c = S.c;
```

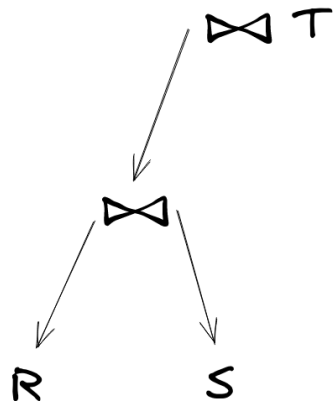
```
{:find [?a ?b ?c]
 :where [[?a :r ?c]
         [?a :t ?b]
         [?b :s ?c]]}
```

# Joins

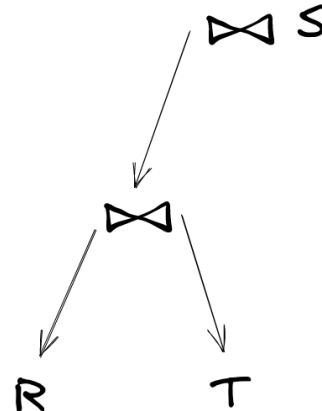
$R(a,c)$              $T(a,b)$              $S(a,c)$

```
SELECT
  *
FROM
  R, T, S
WHERE
  R.a = T.a AND T.b = S.b AND R.c = S.c;
```

```
'{:find [?a ?b ?c]
:where [[?a :r ?c]
        [?a :t ?b]
        [?b :s ?c]]}'
```



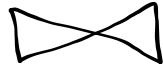
$R \bowtie S \approx 1000$   
 $R \bowtie S \bowtie T \approx 100$



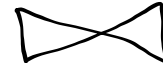
$R \bowtie T \approx 10$   
 $R \bowtie S \bowtie T \approx 100$

# Triangle Query

$G(a,c)$



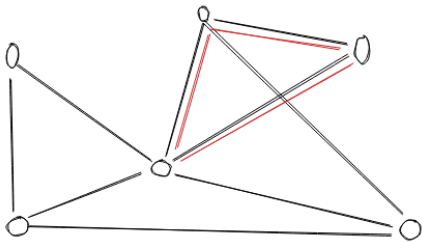
$G(a,b)$



$G(b,c)$

```
SELECT
  g1.f AS a, g1.t AS b, g2.t AS c
FROM
  g AS g1, g AS g2, g AS g3
WHERE
  g1.t = g2.f AND g2.t = g3.t AND g1.f = g3.f;
```

```
{:find [?a ?b ?c]
:where [[?a :g/to ?c]
        [?a :g/to ?b]
        [?b :g/to ?c]]}
```



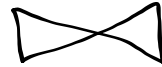
- potentially  $O(N^2)$  rows in an intermediate join
- at most  $O(N^{3/2})$  triangles in any graph

# Worst-case optimal join

A(x)

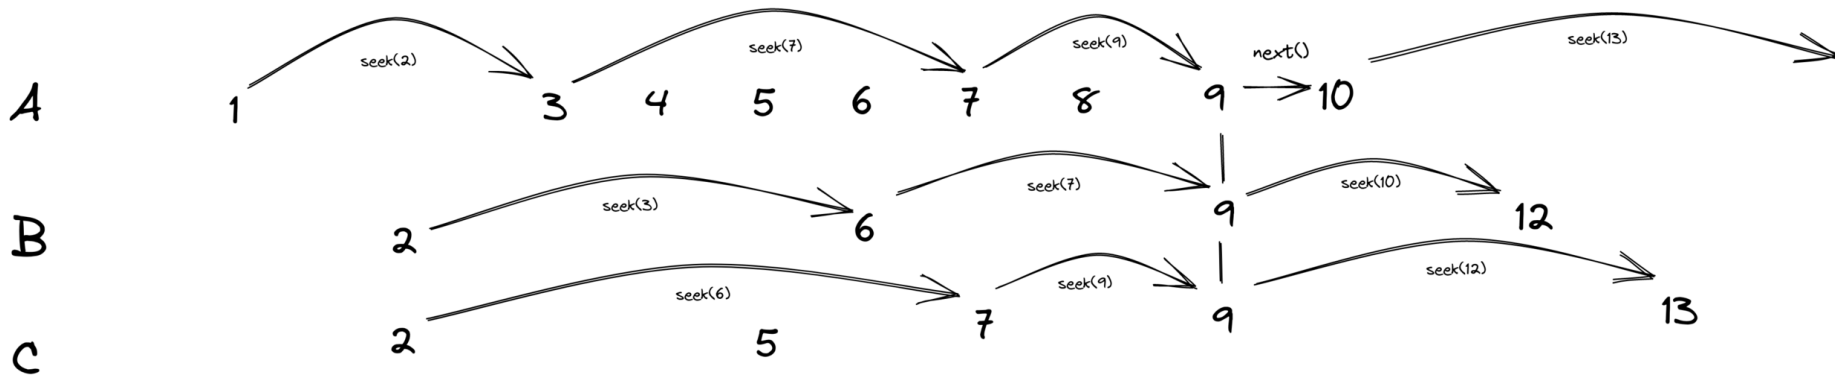


B(x)



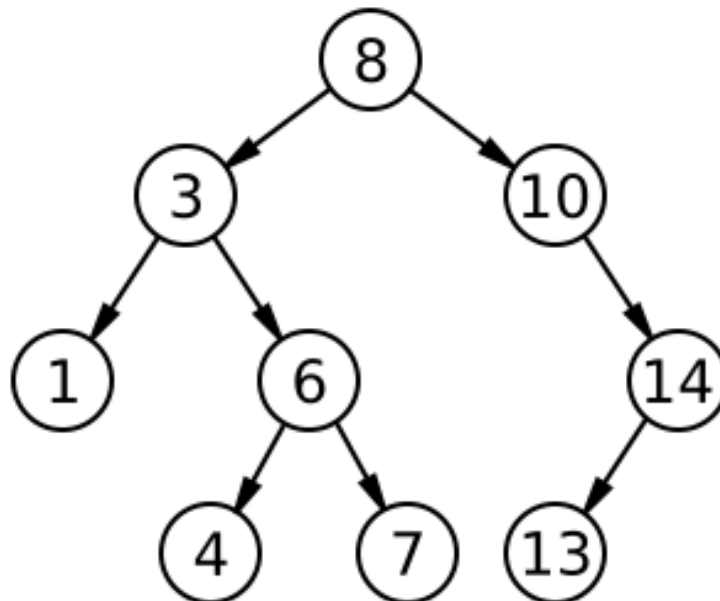
C(x)

```
{:find [?x]
:where [[?x :person/hometown "Berlin"]
        [?x :person/firstname "Bob"]
        [?x :person/eats "Spaghetti"]]}
```



# Unary Iterator

```
(defprotocol LeapIterator  
  (key [this])  
  (next [this])  
  (seek [this k])  
  (at-end? [this]))
```



# Unary Iterator

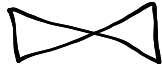
```
(defprotocol LeapIterator
  (key [this])
  (next [this])
  (seek [this k])
  (at-end? [this]))
```

- visiting  $m$  of  $N$  values should have amortized cost  $O(1 + \log(N/m))$
- any `sorted-map` implementation works
- `ISeq` needs to be extended with

```
(defprotocol Seek
  (seek [this k]))
```

# Worst-case optimal join

A(x)



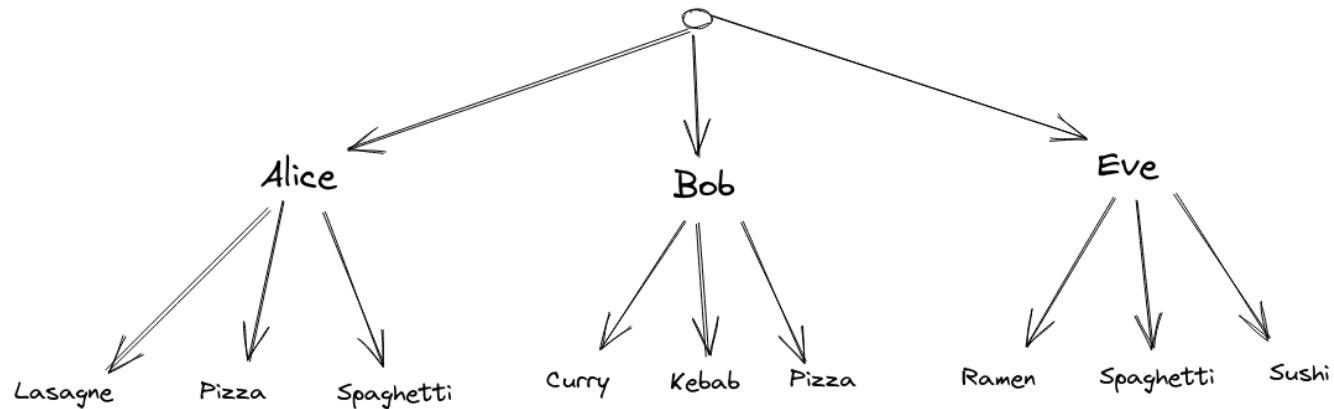
B(x)



C(x,y)

```
{:find [?x ?y]
:where [[?x :person/hometown "Berlin"]
[?x :person/firstname "Bob"]
[?x :person/eats ?y]]}
```

Alice	Lasagne
Alice	Pizza
Alice	Spaghetti
Bob	Curry
Bob	Kebab
Bob	Pizza
Eve	Ramen
Eve	Spaghetti
Eve	Sushi

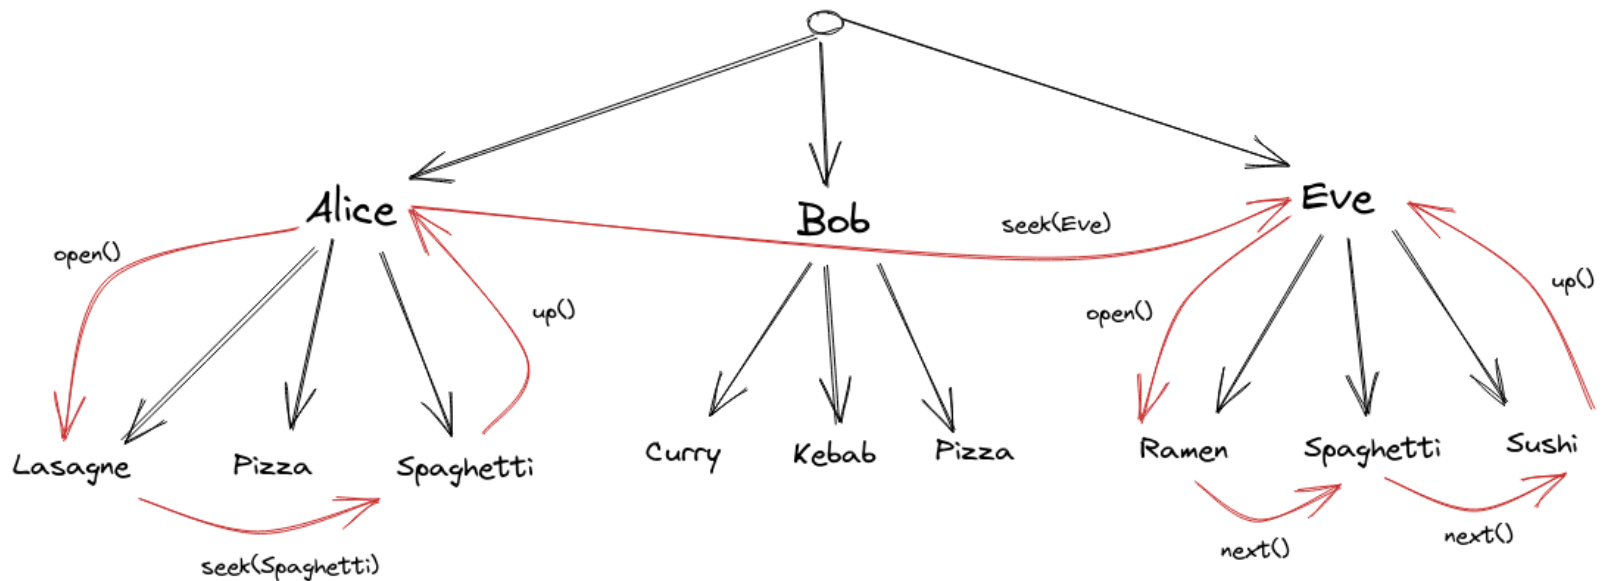




# Multi-arity iterator

```
(defprotocol LeapIterator
  (key [this])
  (next [this])
  (seek [this k])
  (at-end? [this]))
```

```
(defprotocol LeapLevels
  (open [this])
  (up [this])
  (level [this])
  (depth [this]))
```



# A simple datalog engine

```
[ :db/add "Bob" :person/eats "Spaghetti" ]
```

```
(defrecord MemoryGraphIndexed [eav eva ave aev vea vae doc-store opts]
  graph/Graph
  (transact [this tx-data ts] (transact* this tx-data ts))
  ...

  graph/GraphIndex
  (get-iterator [this tuple] (get-iterator* this tuple))
  ...
)
```

# A simple datalog engine

```
[ :db/add "Bob" :person/eats "Spaghetti" ]
```

```
(defn index-triple-add [{opts :opts :as graph} [e a v :as triple]]
  (let [update-in (avl-util/create-update-in avl/sorted-map)
        [he ha hv] (mapv hash triple)]
    (-> graph
      (update-in [:eav he ha] (fnil conj (avl/sorted-set)) hv)
      (update-in [:eva he hv] (fnil conj (avl/sorted-set)) ha)
      (update-in [:ave ha hv] (fnil conj (avl/sorted-set)) he)
      (update-in [:aev ha he] (fnil conj (avl/sorted-set)) hv)
      (update-in [:vea hv he] (fnil conj (avl/sorted-set)) ha)
      (update-in [:vae hv ha] (fnil conj (avl/sorted-set)) he)
      (update :doc-store into [[he e] [ha a] [hv v]]))))
```

# A simple datalog engine

Triple

```
[ :db/add "Bob" :person/eats "Spaghetti" ]
```

"hashed" triple

```
[ :db/add 1 2 3 ]
```

eav index

```
{1 {2 #{3}}}
```

aev index

```
{2 {1 #{3}}}
```

doc-store

```
{1 "Bob" 2 :person/eats 3 "Spaghetti"}
```

# A simple datalog engine

```
{:triple '[?person :person/eats ?food]
:triple-order [:a :e :v]}
```

```
(defn simplify [binding] (map #(if (util/variable? %) '? :v) binding))

(defmulti get-index (fn [graph {:keys [triple] :as _tuple}] (simplify triple)))

(defmethod get-index '[? ? ?]
  [graph {[t1 t2 t3] :triple-order :as _tuple}]
  (get graph (keyword (str (name t1) (name t2) (name t3)))))

(defmethod get-index '[? :v ?]
  [graph {[t1 t2 t3] :triple-order [_ v2 _] :triple :as _tuple}]
  (get-in graph [(keyword (str (name t2) (name t1) (name t3))) v2]))

...
```

# A simple datalog engine

```
(defrecord LeapIteratorAVL [index stack depth max-depth]
  LeapIterator
  (key [this] (first-key index depth max-depth))

  (next [this] (->LeapIteratorAVL (clojure.core/next index) stack depth max-depth))

  (seek [this k]
    (when (seq index)
      (->LeapIteratorAVL (avl/seek index k) stack depth max-depth)))

  (at-end? [this] (empty? index))

  LeapLevels
  (open [this]
    (assert (< (inc depth) max-depth))
    (->LeapIteratorAVL (-> index first second seq) (conj stack index) (inc depth) max-depth))

  (up [this]
    (->LeapIteratorAVL (peek stack) (pop stack) (dec depth) max-depth))

  (level [this] depth)

  (depth [this] max-depth))
```

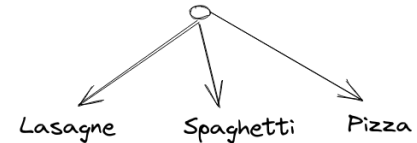
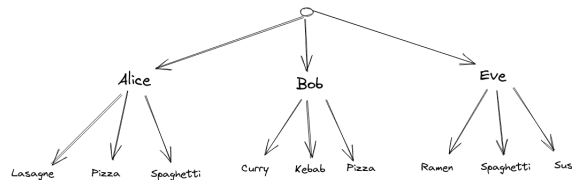
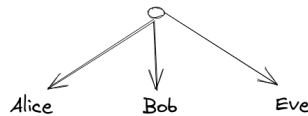
```
(defn ->leap-iterator-avl [index max-depth]
  {:pre [(avl-index? index)]}
  (->LeapIteratorAVL (seq index) [] 0 max-depth))

(->leap-iterator-avl (get-index graph tuple) nb-vars)
```

# A simple datalog engine

```
'{:find [?person ?food]
  :where [[?person :person/hometown "Berlin"]
          [?person :person/eats ?food]
          [?food :food/healthy true]]}'
```

- decide on a variable order ( ?person ?food)
- get iterators



# (simplified) complexity proof

```
for each a in intersect(R, T):  
    for each b in intersect(R[a], S):  
        for each c in intersect(S[b], T[a]):  
            emit((a, b, c))
```

$O(N)$

$O(N)$

$O(N^{3/2})$

- by the proof we emit at most  $O(N^{3/2})$  results
- every intersect is roughly a unary iterator



# Final remarks

- no silver bullet
- variable join order matters
- is it really worth it?
- disk access the heavy part

# References

- Skew strikes back <https://arxiv.org/abs/1310.3314>
- Leapfrog Triejoin <https://arxiv.org/abs/1210.0481>
- Radix Triejoin <https://arxiv.org/abs/1912.12747>
- WCOJ review <https://arxiv.org/abs/1803.09930>
- <https://justinjaffray.com/a-gentle-ish-introduction-to-worst-case-optimal-joins/>
- <http://www.frankmcsherry.org/dataflow/relational/join/2015/04/11/genericjoin.html>
- <https://github.com/FiV0/hooray>